

A Unified Example-Based and Lexicalist Approach to Machine Translation

Davide Turcato Paul McFetridge Fred Popowich and Janine Toole
 Natural Language Laboratory, School of Computing Science, Simon Fraser University
 8888 University Drive, Burnaby, British Columbia, V5A 1S6, Canada
 and
 Gavagai Technology
 P.O. 374, 3495 Cambie Street, Vancouver, British Columbia, V5Z 4R3, Canada
 {turk,mcfet,popowich,toole}@cs.sfu.ca

Abstract

We propose an approach to Machine Translation that combines the ideas and methodologies of the Example-Based and Lexicalist theoretical frameworks. The approach has been implemented in a multilingual Machine Translation system.

1 Introduction

Human translation is a complex intellectual activity and accordingly Machine Translation (henceforth MT) is a complex scientific task, involving virtually every aspect of Natural Language Processing. Many approaches have been proposed, each of them inspired by some insight about translation. Each approach has its own merit, accounting for some aspect of translation better than other approaches, but typically each approach's advantages are countered by weaknesses in other respects. The real challenge is combining different approaches and insights into a comprehensive whole. To this end it is important to compare different approaches, for two reasons:

1. It is important to see to what extent differences are substantial or notational. Sometimes different approaches look at the same subject from different viewpoints, or use different representations, but a formal analysis shows that they are equivalent. This was the case with many formal systems (categorical and phrase structure grammars, finite state machines and regular grammars, explanation-based generalization and partial evaluation, etc.). In other cases differences have been demonstrated to be matters of degree (for instance, in the field of MT, transfer and interlingua approaches).
2. It is important to see to what extent different approaches are mutually exclusive, or whether they can be integrated into one system that encompasses all of them.

In this paper we examine the Example-Based and the Lexicalist approaches to MT. Despite the differences between their paradigms and methods, we argue that:

1. the kind of linguistic resources the two approaches use largely overlap and can be expressed in the same notation;
2. a unified MT architecture can be proposed that encompasses the methodologies of both approaches.

2 The Two Approaches

2.1 Example-Based MT

Arnold et al. (1994:198) outline Example-Based MT (henceforth EBMT) as follows: “The basic idea is to collect a bilingual corpus of translation pairs and then use a best match algorithm to find the closest example to the source phrase in question. This gives a translation template, which can then be filled in by word-for-word translation.”

The paradigm of *translation by analogy* was first introduced by Nagao (1984). In that paper Nagao advocates the use of raw, unanalyzed bilingual data, claiming that linguistic data are more durable than linguistic theories, thus constituting a steadier ground for MT systems. He proposes the use of an unannotated database of examples (possibly collected from a bilingual dictionary) and a set of lexical equivalences simply expressed in terms of word pairs (except for verb equivalences, which are expressed in terms of case frames). The matching process is mainly focused on checking the semantic similarity between the lexical items in the input sentence and the corresponding items in the candidate example.

Many variations and extensions to Nagao’s ideas followed, under different names and acronyms: Example-Based Machine Translation (EBMT, Sumita & Iida 1991), Memory-Based Translation (MBT, Sato & Nagao 1990), Transfer-Driven Machine Translation (TDMT, Furuse & Iida 1992), Case-Based Machine Translation (CBMT, Kitano 1993), etc. There is not always agreement about the usage of such names: for instance, some authors use EBMT and MBT interchangeably, while others keep them distinct. However, all these approaches share the basic idea described above. The main directions in which Nagao’s original model has been extended are the following:

1. Augment the example database with linguistic annotations and, accordingly, perform some linguistic analysis on the input before the matching phase. Sato & Nagao (1990) store examples in the form of pairs of word-dependency trees, along with a set of ‘correspondence links’. For instance, the English-Japanese pair of sentences in (1) is represented as the Prolog facts in (2):

```
(1)  a. He eats vegetables.
      b. Kare ha yasai wo taberu.

(2)  ewd_e([e1,[eat,v],
           [e2,[he,pron]],
           [e3,[vegetable,n]]]).

      jwd_e([j1,[taberu,v],
           [j2,[ha,p],
           [j3,[kare,pron]]],
           [j4,[wo,p],
           [j5,[yasai,n]]])).

      clinks([[e1,j1],[e2,j3],[e3,j5]]).
```

Kitano (1993) proposes the annotation of examples with morphological information for words and then suggests splitting the source and target sentences into *segments*, i.e. continuous sequences of words. He then proposes to annotate examples with a *segment map*, i.e. a correspondence between segments in the source and target sentences, in a similar fashion to what Sato & Nagao (1990) do with word-dependency sub-trees.

2. Explicitly store templates in the bilingual database, instead of sentences (Kaji et al. 1992). A template is a sentence where some phrases have been replaced by variables, annotated with linguistic information. For instance:

(3) X[PRON] eats Y[NP] <-> X[PRON] ha Y[NP] wo taberu

Templates are learnt from pairs of sentences by parsing them, in order to perform correct replacement of words or phrases with annotated variables, and to obtain cross-linguistic variable-sharing.

Furuse & Iida (1992) propose encoding different kinds of bilingual correspondences in the database: string-level correspondences (i.e. plain phrase pairs), pattern-level correspondences (pairs of templates containing variables), grammar-level correspondences (pairs of templates containing variable annotated with syntactic categories, like those proposed by Kaji et al.). Moreover, they associate a source expression with several target expressions, each of which is provided with a set of examples that show the contexts in which each target expression can be correctly used. For instance:

(4) X wo o-negaishimasu <->
 may I speak to X' ((jimukyoku{office}), ...)
 please give me X' ((bangou{number}), ...)

where X' is the translation of X and each (X{X'}) pair in parentheses is an instantiation of such a translation pair.

3. Obtain the translation of a complete sentence by utilizing more than one translation example and combine some fragments of them. For instance, Sato & Nagao (1990) show how to obtain the translation (5), given the examples (6) and (7).

(5) *He buys a book on international politics* ↔
 Kare ha kokusaiseiji nitsuite kakareta hon wo kau

(6) **He buys** a notebook ↔
 Kare ha nouto wo kau

(7) *I read a book on international politics* ↔
 Watashi ha **kokusaiseiji nitsuite kakareta hon** wo yomu

For an input sentence, they construct a *matching expression*, i.e. a pointer to a *translation unit*. A translation unit is a word-dependency sub-tree to be found in the example database (the $e_1, \dots, e_n, j_1, \dots, j_m$ of example

(2)). Such pointers can be optionally followed by a list of commands for deletion/replacement/adjunctions of nodes dominated by the node pointed to. The replaced or adjoined elements are other matching expressions. For instance, given (2), a matching expression for the sentence (8) might be (9).

(8) *He eats mashed potatoes*

(9) $[e1, [r, e3, [e_x]]]$

which represents the tree obtained by replacing (r for ‘replace’) $e3$ with e_x in $e1$. In turn, e_x is a pointer to a sub-tree for *mashed potatoes* in some other example.

As several matching expressions can be candidates for the same input sentence, Sato & Nagao define a scoring system for competing translation units, based on their length (the longer, the better) and the semantic similarity between their contexts, i.e. the input sentence and the example from which the translation unit is taken (the more similar, the better).

2.2 Lexicalist MT

Lexicalist MT (henceforth LMT) is a variant of the transfer approach to MT. In LMT transfer is a mapping between bags of lexical items, instead of trees (Whitelock 1994).

The first step of the translation process is the analysis of an input sentence. Such analysis is performed on a purely monolingual basis, independently from considerations of translation direction and language pair. The same kind of declarative grammars are used for parsing and generation. Moreover, grammars tend to follow a lexicalist, sign-based¹ approach (Pollard & Sag 1994). Grammar rules are reduced to a small number of general rule schemata. Lexical items are multidimensional signs containing all the information about their modes of combination (subcategorization, head-modifier relations, etc.). As a result of parsing, lexical items are instantiated with indices expressing their interdependencies with other lexical items in the sentence.

Transfer is a mapping from a bag of instantiated source lexical items resulting from parsing to a corresponding bag of target lexical items. Bilingual knowledge is reduced to a bilingual lexicon, augmented with cross-linguistic correspondences in the form of equated variables. Transfer is performed by finding a set of bilingual entries that covers the source bag. The target bag is comprised of the target sides of the selected bilingual entries. Phrasal and idiomatic expressions are accounted for by multi-word bilingual entries, where lexical items on either side have no inherent order and can be discontinuous in the input or output sentence. A schematic bilingual entry is shown in (10), where subscripts represent indices encoding word dependencies.

(10) $eat : v_{a,b,c} \leftrightarrow taberu : v_{a,d,e} \ \& \ ha : p_{d,b} \ \& \ wo : p_{e,c}$

Generation orders the target lexical items into a grammatical sentence, according to a target grammar and to the constraints expressed by the indices instantiated on target lexical items as a result of transfer.

¹Following the Saussurean approach taken in HPSG, we define signs as “structured complexes of phonological, syntactic, semantic, discourse and phrase-structural information” (Pollard & Sag 1994:15).

2.3 Comparison

It is interesting to note that the introduction of both EBMT and LMT was motivated by the rejection of structural transfer, due to its inadequacy to cope with structurally divergent languages like English and Japanese (Nagao 1984, 179; Whitelock 1994, 343–345). However, the two approaches differ in the way they avoid the recursive traversal of an analysis tree structure in transfer.

In EBMT structural transfer is avoided by adopting the following guidelines:

1. Sentence-level correspondences are covered via an explicit stipulation of all such possible correspondences in the bilingual knowledge base. Therefore EBMT advocates a bilingual knowledge base stating equivalences between the maximal translation units, i.e. sentences (or sentence templates).
2. Given such flat structure of the bilingual database, no deep linguistic analysis is required. Transfer is performed by looking up the bilingual database for a suitable match to the input sentence. Linguistic analysis is only performed to the extent it is necessary to effectively perform the template matching.

In LMT the following guidelines are adopted:

1. A full linguistic analysis is performed on input sentences. As a result of the lexicalist, sign-based approach to parsing and the use of indices to represent dependencies among lexical items, individual lexical items contain all the information about their structural relationships with the other lexical items.
2. Given that all the information about a sentence structure is stored in lexical items, transfer can be reduced to a mapping of a bag of source lexical items onto a bag of target lexical items. Therefore LMT advocates a bilingual knowledge base stating equivalences between the minimal translation units, i.e. lexical items. Information about word order is also dropped from transfer, as it is considered a monolingual issue, accounted for by the linear precedence constraints expressed in grammar rules. The dependencies expressed by indices are the only information that must be necessarily transferred from source to target lexical items.

3 A Unified Bilingual Knowledge Base

Despite the different and somehow antithetic architectures, the kind of bilingual resources required by the two approaches tend to converge. We show that it is possible to define a bilingual knowledge base in such a way that it can serve both approaches.

At a formal level, it can be shown that the kind of information used by the two approaches largely overlaps. The information needed for EBMT systems can be adequately expressed in a LMT notation. We list here some parallelisms:

1. Case frames as used in EBMT correspond to subcategorization frames in LMT. Therefore an EBMT case frame is equivalent to an LMT verb lexical entry. Aside from word order issues, which will be dealt with later, the same holds for templates where some arguments are left unspecified, as a comparison between (3) and (10) shows.

2. A comparison between (2) and (10) shows that Sato & Nagao's (1990) word-dependency trees contain the same information as LMT bilingual entries: words, grammatical descriptions, monolingual dependencies, cross-linguistic correspondences.
3. TDMT templates correspond to sets of bilingual entries.

A bilingual lexicon as used in LMT can adequately represent all the information needed in EBMT. Therefore, we advocate the use of the LMT notation as a theory neutral knowledge representation language that can equally support EBMT and LMT bilingual knowledge bases. The adoption of such notation does not commit one to using one or the other approach. Such a bilingual resource is close in spirit to the kind of Bilingual Knowledge Bank advocated by Sadler & Vendelmans (1990).

The neutrality of the proposed notation also relies on the fact the notation's semantics is underspecified. The notation is such that it can be interpreted in different ways, in developing and using a knowledge base. Particularly, a bilingual entry's source or target side can be interpreted as either a bag or a sequence. In the latter case, word order is relevant in matching some input with a bilingual entry. A further constraint on the matching procedure may be the requirement that input words matching bilingual entry items must be contiguous in the sentence. If both order and contiguity constraints are activated, then a bilingual knowledge base is interpreted as a knowledge base of sentences (or phrases), as in Nagao's (1984) original proposal, or segments, as proposed by Kitano (1993). If the order constraint is activated and the contiguity constraint is dropped, then bilingual entries represent templates. If both constraints are dropped, then bilingual entries represent word-dependency trees or LMT lexical bags. Therefore the same notation can be used with different ideas in mind and, to some extent, the same knowledge base can be reused under different interpretations.

Our experience in large scale bilingual lexical development (English-Spanish) showed that the commitment to a specific semantics may even be changed after a bilingual knowledge base has been developed, if some conventions in writing entries are observed. We remarked that our lexicographers spontaneously used the obvious convention of writing bilingual entry items in the same order in which words appear in sentences. With some exceptions (for instance, Spanish verbs accompanied by clitic pronouns), the order in which bilingual entry items can appear in sentences turned out to be unique in most cases. This gave us the choice of interpreting our bilingual entries as either bags or sequences, which was an option unforeseen at the beginning. It is also possible to choose a mixed semantics, e.g. interpreting source sides as sequences and target sides as bags. Different considerations come into play for different languages and translation directions. For instance, the order constraint might be appropriate for a language with a relatively fixed word order, but not for one with a relatively free word order (even more so, when the language at hand is used as a source language).

The notation can also be extended to contain explicit place-holders for missing elements, thus resembling templates more closely. For an illustration of such an extension see (Turcato et al. 1997).

Besides notation, a second issue is the actual information that the two approaches require of a bilingual knowledge base. As noted above, EBMT tends to require equiv-

alences between maximal translation units, i.e. sentences, while LMT tends to require equivalences between minimal units, i.e. lexical items. Such a divergence is actually less dramatic if we take a closer look at the issue. To this end we introduce a distinction between two different purposes of an example database.

1. Examples provide information about sentence well-formedness. The required amount of such information is inversely proportional to the amount of linguistic analysis performed on the input. At one extreme, we have the case that no analysis is performed. In this case, all the possible sentences should be listed in the bilingual database. The next level is when input words are assigned syntactic categories. In this case, sentences in the bilingual database can be either replaced by or used as templates. At the opposite extreme is the case where a complete linguistic analysis is performed. In this case, examples are no longer needed to provide information about well-formedness. Lexical equivalences are sufficient. For instance, if we assume that a sentence input is analyzed into a word-dependency tree, then a bilingual example like (2) can be replaced by a set of bilingual lexical entries without any loss of information (provided, as is assumed by a lexicalist approach, that each lexical item contains information about the arguments it subcategorizes for). Therefore a lexicalist approach can be seen as the lower bound on a continuous scale of different EBMT approaches, depending on the amount of linguistic analysis performed.
2. Examples provide information about non-compositional translations (e.g. idioms) and contrastive information about different ways in which a word is translated in different contexts (sense-ambiguous words). This is the case of examples like (4), for instance. This kind of information is equally required by EBMT and LMT systems, regardless of the chosen approach to linguistic analysis, and needs to be expressed in either case by multi-word bilingual entries.

To sum up, the required information is the same in EBMT and LMT, to some extent. The extent of the residual difference is a matter of degree of linguistic analysis performed by the system.

4 A Unified Architecture

Arnold et al. (1994:201) suggest that “there is no radical incompatibility between example-based and rule-based approaches, so that the challenge lies in finding the best combination of techniques from each. Here one obvious possibility is to use traditional rule-based transfer as a fall back, to be used only if there is no complete example-based translation.” Rather than proposing a multi-engine approach with a duplication of resources, we propose a single architecture that encompasses the two approaches and integrates the basic tenets of both.

A common characteristic of all EBMT approaches is that the translation process is driven by the content of the bilingual knowledge base. The core operation of all such approaches is the match of an input sentence against examples. It is the result of such a match that drives further computation, in terms of calculating similarity,

replacing items in the chosen example, combining fragments of different examples (this prioritization of transfer is made explicit in approaches like TDMT). On the contrary, in LMT different translation steps are clearly separated. As pointed out, parsing is performed on a purely monolingual ground, regardless of the specific translation flow in which it occurs. We propose a translation architecture that combines the advantages of the two approaches, by using bilingual information to drive the translation process, while preserving the modularity of the system.

A bilingual knowledge base as described above is not only a source of bilingual information, but it also encodes a considerable amount of monolingual linguistic information, on either side. A multi-word bilingual entry gives syntactic and semantic information about the analysis of phrasal expressions, collocations and idioms. Even single-word entries give clues about the analysis of lexically ambiguous items.

We propose to take advantage of this source of information to drive the parsing process of an input sentence. Given the search space defined by the monolingual lexicon and grammar, the information contained in the bilingual lexicon is used to prioritize certain analyses over others². A bilingual lexicon lookup before parsing offers a partial analysis (in terms of lexical disambiguation, dependencies and, optionally, word order), which is tried before any other hypothesis supported by the monolingual lexicon and grammar. If, for instance, chart parsing is in use, the example-based approach amounts to prioritizing edges in the chart agenda. Moreover, edges licensed by the same multi-word bilingual entry are assigned a common identifier, so as to ensure that they all fail or succeed together.

When several bilingual entries apply, they are prioritized by the cardinality of their side being used. This sorting mechanism implements a kind of ‘elsewhere condition’: more specific entries override more general ones. This device can be regarded as a lexicalist implementation of the scoring mechanism described by Sato & Nagao (1990), according to which longer translation units are preferred over shorter ones.

We show how the translation process works with an English-Spanish example:

(11) *They cut back on investments*

Let’s assume that the English lexicon and the bilingual lexicon contain, respectively, the following (simplified) entries:

back : adv_a
back : n_a
cut : $\text{iv}_{a,b}$
cut : $\text{iv}_{a,b,c}$
cut : $\text{tv}_{a,b,c}$
investment : n_a
on : $\text{p}_{a,b}$

²A similar idea has been proposed by Kinoshita (1998:80) in a different theoretical framework.

a.	<i>back</i> :adv _a	↔ <i>atrás</i> :adv _a
b.	<i>back</i> :n _a	↔ <i>espalda</i> :n _a
c.	<i>cut</i> :iv _{a,b}	↔ <i>cortar</i> :iv _{a,b}
d.	<i>cut</i> :iv _{a,b,c} & <i>back</i> :adv _c	↔ <i>hacer</i> :tv _{a,b,d} & <i>economía</i> :n _d
e.	<i>cut</i> :iv _{a,b,c} & <i>back</i> :adv _c & <i>on</i> :p _{a,d}	↔ <i>reducir</i> :tv _{a,b,d}
f.	<i>cut</i> :tv _{a,b,c}	↔ <i>cortar</i> :tv _{a,b,c}
g.	<i>investment</i> :n _a	↔ <i>inversión</i> :n _a
h.	<i>on</i> :p _{a,b}	↔ <i>en</i> :p _{a,b}

Given such monolingual and bilingual lexical entries, we show below all the possible ways in which the input sentence (11) can be grammatically covered in parsing and correspondingly translated (we omit details about *they*, which is syntactically unambiguous and is dropped in Spanish). The solutions are listed in no specific order. Note that more than one translation can be given for the same parse, depending on what bilingual entries are used. Conversely, different parses can result in the same translation. At the end of each line we also indicate what bilingual entries have been used.

<i>cut</i>	<i>back</i>	<i>on</i>	<i>investments</i>		
iv _{a,b}	adv _a	p _{a,c}	n _c	→ <i>cortan atrás en las inversiones</i>	{cahg}
tv _{a,b,c}	n _c	p _{a,d}	n _d	→ <i>cortan espalda en las inversiones</i>	{fbhg}
tv _{a,b,c}	n _c	p _{c,d}	n _d	→ <i>cortan espalda en las inversiones</i>	{fbhg}
iv _{a,b,c}	adv _c	p _{a,d}	n _d	→ <i>hacen economías en las inversiones</i>	{dhg}
iv _{a,b,c}	adv _c	p _{a,d}	n _d	→ <i>reducen las inversiones</i>	{eg}

Note that in our specific example it is irrelevant whether the order and contiguity constraints are enforced on the bilingual lexicon. The parsing strategy we propose first tries to find a parse consistent with the source side of bilingual entry (e), the longest available. Therefore, assuming that no failure occurs, the first translation returned is *reducen las inversiones*, which is the most correct. If a failure occurs anywhere down the path for all the parses covered by the source side of bilingual entry (e), the source side of bilingual entry (d) is tried next. Therefore, *hacen economías en las inversiones* would be the second translation returned. When the bilingual lexicon offers no way of prioritizing among parsing hypotheses, any other available prioritization mechanism can still be used. Also, in using a bilingual lexicon, different strategies could be used. For instance, an alternative to using the longest match first, would be to look for the cover that uses the fewest number of bilingual entries.

As discussed above, Sato & Nagao (1990) also use a second scoring mechanism, based on similarity between the input sentence and the candidate translation units. A lexicalist counterpart of such a mechanism would amount to a word sense disambiguation module, provided that senses are associated with words in the bilingual lexicon. In fact, the problem of choosing the right translation for a word or phrase that can be translated in different ways amounts to choosing the correct word sense for that word or phrase. This, in turn, is customarily done in the word sense disambiguation literature by looking at the context in which the word or phrase occurs (e.g. Resnik 1995; Yarowsky 1995), thus paralleling Sato & Nagao's (1990) idea. Although we have not

implemented any such word sense disambiguation module, it would be straightforward to incorporate such a module in the system architecture without affecting the other modules. As for associating senses with lexical items in the bilingual lexicon, a method for automatically selecting and ranking bilingual entries (unannotated for sense), based on an input word's sense and context, has been proposed by Sanfilippo & Steinberger (1997).

Besides the monolingual considerations discussed above, this approach also has the advantage of biasing parsing towards analyses that are supported by the bilingual lexicon. An analysis, even a correct one, is useless if the transfer component does not have the means to map it onto a target representation. Therefore it is a practical choice to prioritize those analyses that are amenable to a successful transfer.

The proposed approach, which has been incorporated into a large scale MT system (Popowich et al. 1997), does not affect the results provided by the system, it only affects the order in which they are provided. Of course, if a system returns the first solution found, the system behavior indeed changes.

5 Conclusion

A knowledge representation format and a system architecture have been proposed that allow an effective integration of Example-Based and Lexicalist approaches to MT into a unified approach, which we call Example-Based Lexicalist Machine Translation (EBLMT). This approach combines the advantages of each approach. From the point of view of LMT, it uses bilingual knowledge to drive parsing, providing additional information to solve syntactic ambiguities and prioritizing the parsing agenda in a more efficient way. From the point of view of an EBLMT system like (Sato & Nagao 1990), for instance, it allows the removal of the bilingual database's redundancy coming from the overlap of examples. Moreover, the flexibility of the knowledge representation format and the modularity of the architecture allow a system to work in different modalities, by simply setting some system parameters.

Acknowledgements

This research was supported by a Collaborative Research and Development Grant from the Natural Sciences and Engineering Research Council of Canada (NSERC), and by the BC Advanced Systems Institute (ASI).

References

- Arnold, D., L. Balkan, R. Lee Humphreys, S. Meijer & L. Sadler: 1994, *Machine Translation. An Introductory Guide*, Manchester-Oxford: NCC Blackwell.
- Furuse, Osamu & Hitoshi Iida: 1992, 'Cooperation between transfer and analysis in example-based framework', in *Proceedings of the Fifteenth [sic] International Conference on Computational Linguistics (COLING-92)*, Nantes, France, pp. 645–651.
- Kaji, Hiroyuki, Yuuko Kida & Yasutsugu Morimoto: 1992, 'Learning translation templates from bilingual text', in *Proceedings of the Fifteenth [sic] International Conference on Computational Linguistics (COLING-92)*, Nantes, France, pp. 672–678.

- Kinoshita, Jorge: 1998, 'A transfer dictionary for words and bigrams', in Flávio Moreira de Oliveira, ed., *Advances in Artificial Intelligence — 14th Brazilian Symposium on Artificial Intelligence, SBIA'98, Porto Alegre, Brazil, November 4–6, 1998 — Proceedings*, Berlin: Springer, pp. 73–82.
- Kitano, Hiroaki: 1993, 'A comprehensive and practical model of memory-based Machine Translation', in *Proceedings of the 13th International Joint Conference on Artificial Intelligence*, Chambery, France, pp. 1276–1282.
- Nagao, Makoto: 1984, 'A framework of a mechanical translation between Japanese and English by analogy principle', in Alick Elithorn & Ranan Banerji, eds., *Artificial and Human Intelligence: edited review papers at the International NATO Symposium on Artificial and Human Intelligence sponsored by the Special Programme Panel held in Lyon, France October, 1981*, Amsterdam, North-Holland: Elsevier Science Publishers, chap. 11, pp. 173–180.
- Pollard, Carl & Ivan Sag: 1994, *Head-driven Phrase Structure Grammar*, Stanford University, CA: Centre for the Study of Language and Information.
- Popowich, Fred, Davide Turcato, Olivier Laurens, Paul McFetridge, J. Devlan Nicholson, Patrick McGivern, Maricela Corzo-Pena, Lisa Pidruchney & Scott MacDonald: 1997, 'A lexicalist approach to the translation of colloquial text', in *Proceedings of the 7th International Conference on Theoretical and Methodological Issues in Machine Translation*, Santa Fe, New Mexico, USA, pp. 76–86.
- Resnik, Philip: 1995, 'Disambiguating noun groupings with respect to Wordnet senses', in *Proceedings of the Third Workshop on Very Large Corpora*, Cambridge, Massachusetts, USA, pp. 54–68.
- Sadler, Victor & Ronald Vendelmans: 1990, 'Pilot implementation of a bilingual knowledge bank', in *Proceedings of the 13th International Conference on Computational Linguistics (COLING-90)*, Helsinki, Finland, pp. 449–451.
- Sanfilippo, Antonio & Ralf Steinberger: 1997, 'Automatic selection and ranking of translation candidates', in *Proceedings of the 7th International Conference on Theoretical and Methodological Issues in Machine Translation*, Santa Fe, New Mexico, USA, pp. 200–207.
- Sato, Satoshi & Makoto Nagao: 1990, 'Toward memory-based translation', in *Proceedings of the 13th International Conference on Computational Linguistics (COLING-90)*, Helsinki, Finland, vol. 3, pp. 247–252.
- Sumita, Eiichiro & Hitoshi Iida: 1991, 'Experiments and prospects of example-based Machine Translation', in *Proceedings of the 29th Annual Meeting of the Association for Computational Linguistics (ACL-91)*, Berkeley, CA, USA, pp. 185–192.
- Turcato, Davide, Olivier Laurens, Paul McFetridge & Fred Popowich: 1997, 'Inflectional information in transfer for lexicalist MT', in *Proceedings of the International Conference 'Recent Advances in Natural Language Processing' (RANLP-97)*, Tzigov Chark, Bulgaria, pp. 98–103.
- Whitelock, Pete: 1994, 'Shake and bake translation', in C.J. Rupp, M.A. Rosner & R.L. Johnson, eds., *Constraints, Language and Computation*, London: Academic Press, pp. 339–359.
- Yarowsky, David: 1995, 'Unsupervised word sense disambiguation rivaling supervised methods', in *Proceedings of the 33th Annual Meeting of the Association for Computational Linguistics (ACL-95)*, Cambridge, Massachusetts, USA, pp. 189–196.